

AYUSH RANJAN ROY

AI Engineer | Multi-Agent Systems, RAG, LLM Infrastructure & Cloud Platform

Gurugram, India | royayush0611@gmail.com | +91-9540968483 | linkedin.com/in/ayush-ranjan-roy | github.com/AyushRanjanRoy-01

Claude Certified Architect - Foundations (Anthropic, May 2026) | ayushranjanroy.com

SUMMARY

AI Engineer shipping production multi-agent systems and RAG pipelines (LangGraph, Claude, Azure AI, MCP, HITL), backed by 2 years modernizing Accenture MxDR / Adaptive MxDR (AMxDR) AWS cloud architecture (Terraform, secrets, secure telemetry, CI/CD). Delivered multi-agent O2C document intelligence from PoC to 300+ docs/day at 97%+ accuracy; cut multi-page inference latency ~55% and lifted classification accuracy ~40% via ensemble routing and async parallelism.

EXPERIENCE

Full Stack LLM Development Analyst / AI Engineer

Aug 2025 - Present

Accenture - GenAI & Data, Gurugram | Multi-Agent Systems & LLM Infrastructure

- Built multi-agent systems for enterprise Order-to-Cash / Dispute & Deduction domains using LangGraph orchestration, durable workflow state, and human-in-the-loop control points.
- Orchestrated 4 LLMs (Claude Sonnet, GPT-4o/5, Azure Document Intelligence) for classification & extraction; scaled from PoC to 300+ documents/day at 97%+ accuracy with ensemble majority-voting and a hot-reload model gateway.
- Designed async parallel inference with confidence-score routing and per-model cost control - lifted classification accuracy ~40%, scaled batch capacity 9 to 50 pages, cut latency ~55% multi-page / ~30% single-page.
- Implemented MCP-style governed tool access (e.g. DB tools) and operator-facing agent observability with OpenTelemetry + Datadog; hardened multi-agent runs across 9 failure scenarios with zero unhandled errors in production.
- Authored architecture docs and ran LangGraph/LangChain KT for 4 engineers - cut ramp-up time ~60%.

Platform Engineer

Aug 2023 - Aug 2025

Accenture - Security | MxDR / Adaptive MxDR (AMxDR) AWS Cloud Platform

- Drove MxDR to Adaptive MxDR (AMxDR) platform AWS cloud modernization as lead engineer: Terraform IaC, GitHub Actions CI/CD, and live-incident debugging across ~20 microservices serving 800+ enterprise accounts - cut deployment turnaround ~35% and infra cost ~30%.
- Owned Terraform and server management across 6 environments; modernized 8+ AWS Lambda functions (Python 3.8 to 3.12) cutting runtime failures ~50%; multi-env blue-green (Route 53, ASG) with zero customer-facing downtime during a live pipeline upgrade.
- Automated AWS Secrets Manager migrations + IAM updates; modernized 12+ Java/Python services; restricted-egress proxy automation; CodeQL/Dependabot remediation and runtime hardening.
- Deployed mTLS / OpenTelemetry gRPC synthetic health (5-min checks) into enterprise SIEM paths; PII log redaction and multi-tenant production incident support (Datadog/CloudWatch).

OPEN SOURCE PROJECTS

HelixOps - AI-SRE Incident IDE | github.com/AyushRanjanRoy-01/HelixOps

- 5-agent RCA pipeline with RAG over runbooks, confidence-scored diagnoses, human-gated remediation, JWT/RBAC, mock-safe defaults, full decision-trace audit logging.

RAGGym - RAG Learning Platform | rag-gym.streamlit.app

- Swappable RAG stack (LLM, embeddings, vector store), dense + RRF retrieval, multi-query, reranking, self-correction loop, RAGAS evaluation harness, AI-graded practice mode; Llama 3.2 local via Ollama or API models.

TECHNICAL SKILLS

Agents & GenAI: LangGraph, LangChain, Multi-Agent Orchestration, MCP, HITL, RAG, FSM checkpoints, Prompting

Models & Cloud AI: Claude, GPT-4o/5, Azure AI Foundry / Document Intelligence, ensemble routing, RAGAS, cost/latency telemetry

Data: PostgreSQL, SQL, CTEs, Redis/ElastiCache, pgvector, Qdrant, async Python, FastAPI

Platform & Security: AWS (Lambda, EC2, S3, ASG, Route 53), Terraform, Secrets Manager, IAM, mTLS, SIEM telemetry, Docker, GHA, CodeQL

Observability: OpenTelemetry, Datadog, CloudWatch, structured tracing, synthetic health checks

EDUCATION

B.Tech, Electronics and Communication Engineering - VIT Chennai

2019 - 2023

Coursework: Artificial Intelligence; Essentials of Machine Learning; Cryptography and Network Security; Information Security Analysis and Audit; Data Structures and Algorithms; Operating Systems.