

AYUSH RANJAN ROY

AI Engineer | Multi-Agent Systems, RAG, Platform & Cloud Security Foundations

Gurugram, India | royayush0611@gmail.com | +91-9540968483

linkedin.com/in/ayush-ranjan-roy | github.com/AyushRanjanRoy-01 | ayushranjanroy.com

SUMMARY

AI Engineer building production multi-agent systems and RAG-backed applications, with a foundation in cloud platform engineering and security controls (IaC, secrets, secure telemetry, CI/CD). Focus: LangGraph orchestration, human-in-the-loop gates, MCP-style tool access, observability, and reliable cloud infrastructure.

EXPERIENCE

Full Stack LLM Development Analyst / AI Engineer

Aug 2025 - Present

Accenture - GenAI & Data, Gurugram

- Build multi-agent systems for enterprise Order-to-Cash domains, including Dispute & Deduction, using LangGraph orchestration, durable workflow state, and human-in-the-loop control points.
- Design async multi-model inference services with retries, backoff, concurrency control, and latency/token/cost telemetry.
- Implement MCP-style governed tool access and operator-facing agent observability with OpenTelemetry and Datadog tracing.
- Contribute RAG-style retrieval and knowledge grounding alongside agent decision paths.
- Author technical documentation and run knowledge-transfer sessions on LangGraph/LangChain patterns.

Platform Engineer

Aug 2023 - Aug 2025

Accenture - Security / Cloud Platform, Gurugram

- Modernized cloud infrastructure with Terraform, multi-environment CI/CD, and Jenkins to GitHub Actions migrations.
- Owned Redis/ElastiCache-backed shared metadata platform services, certificates, Auto Scaling, and environment consistency.
- Automated AWS Secrets Manager migrations and IAM updates; remediated CodeQL/Dependabot findings; hardened container/runtime baselines.
- Migrated Lambda and microservice runtimes to supported Python/OS versions; reduced unused compute footprint.
- Built restricted-egress proxy automation and synthetic health monitoring over mTLS/gRPC OpenTelemetry into enterprise SIEM paths; supported multi-tenant production incidents and blue-green deployments.

OPEN SOURCE PROJECTS

HelixOps - AI-SRE Incident IDE

- Multi-agent RCA with RAG over runbooks, human-approved remediation, JWT/RBAC, mock-safe defaults.
github.com/AyushRanjanRoy-01/HelixOps

RAGGym - RAG Learning Platform

- Swappable LLM/embeddings/vector store, self-correcting retrieval, RAGAS evaluation, practice mode. rag-gym.streamlit.app
| github.com/AyushRanjanRoy-01/RaGGym-Chatbot-CLI

TECHNICAL SKILLS

Agents & GenAI: LangGraph, LangChain, Multi-Agent Orchestration, MCP, HITL, RAG, Prompting

Models & Eval: Claude, GPT-4o/5, Azure AI Foundry, RAGAS, ensemble/confidence routing

Data: PostgreSQL, SQL, pgvector, Qdrant, async Python, FastAPI

Security & Platform: Terraform, AWS (Lambda, EC2, S3, ASG), Secrets Manager, IAM, mTLS, SIEM telemetry, Docker, GitHub Actions, CodeQL/Dependabot

Observability: OpenTelemetry, Datadog, CloudWatch, structured tracing

EDUCATION

B.Tech, Electronics and Communication Engineering

Vellore Institute of Technology, Chennai | 2019 - 2023
Relevant coursework: Artificial Intelligence; Essentials of Machine Learning; Cryptography and Network Security; Information Security Analysis and Audit; Data Structures and Algorithms; Operating Systems.

CERTIFICATION

Claude Certified Architect - Foundations | Anthropic | May 2026